

# Introduction to TensorFlow

Tinghui Wang

# Deep Learning in One Slide



## What is it?

Extract useful patterns  
(features) from data



## How

Neural Network  
Optimization



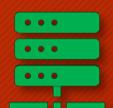
## How (Practical)

Python  
Tensorflow  
Friends!



## Hard Part

Good questions  
Good Data



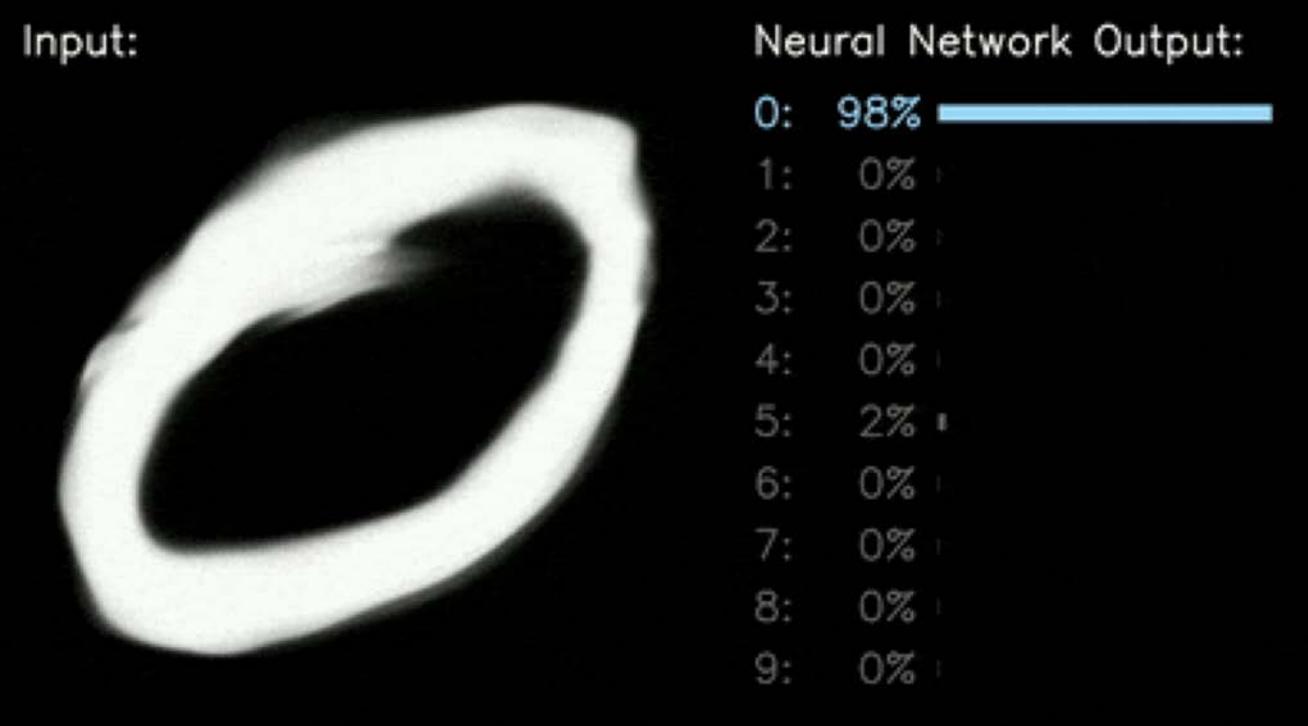
## Why now?

Data  
Hardware  
Community  
Tools  
Investment

# Deep Learning: Exciting Progress

- **Exciting progress:**
  - Face recognition
  - Image classification
  - Speech recognition
  - Text-to-speech generation
  - Handwriting transcription
  - Machine translation
  - Medical Diagnosis
  - Digital assistants
  - Ads, search, social recommendation
  - Game playing with deep RL
  - Auto-drive

## Handwriting Digit Classification

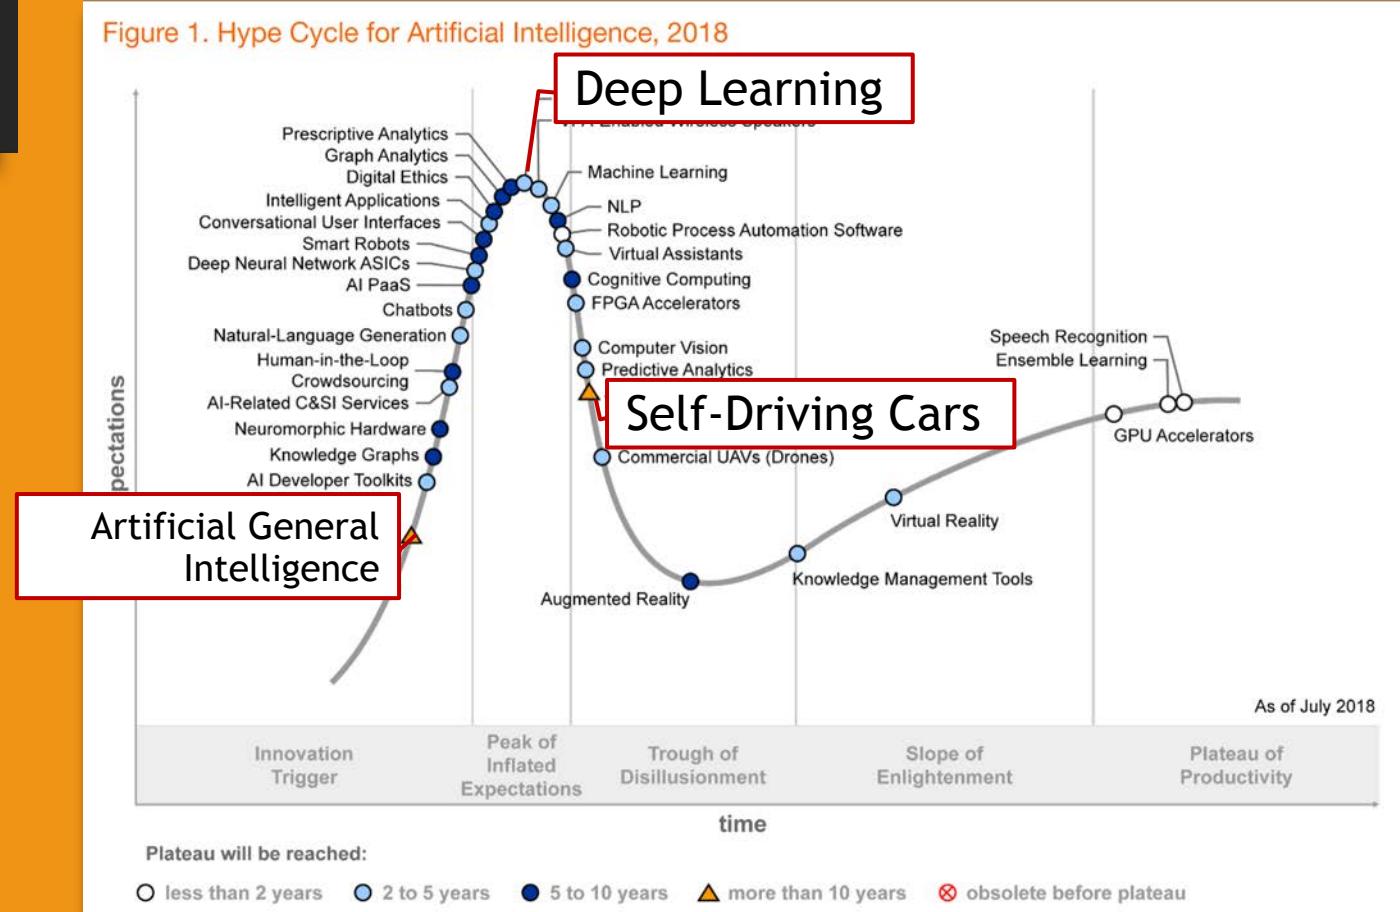


From: Deep Learning Basics, MIT (<https://github.com/lexfridman/mit-deep-learning>)

# Hype Cycle: Artificial Intelligence

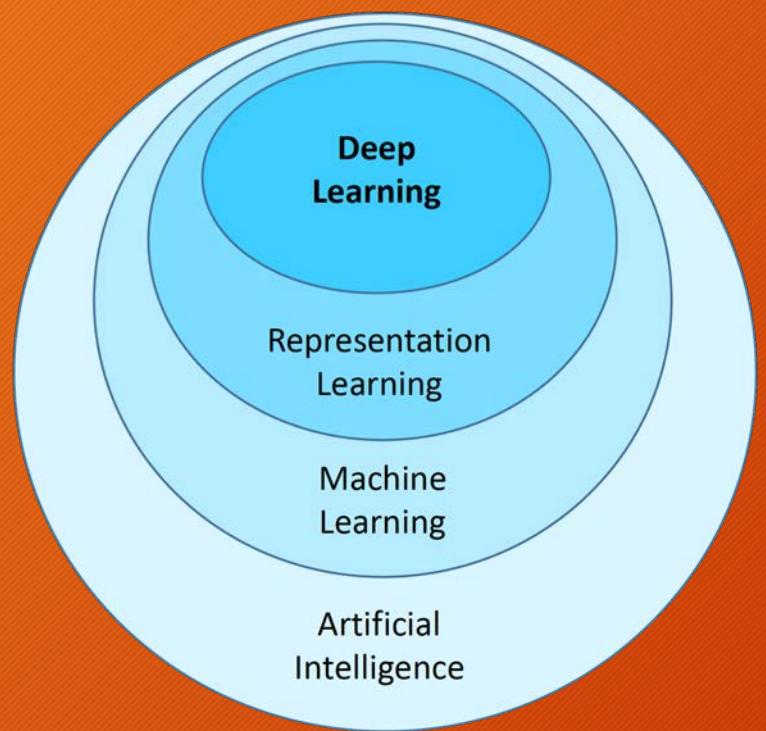
*"There's not going to be an AI winter, because it drives your cellphone. In the old AI winters, AI wasn't actually part of your everyday life. Now it is."*

-- Geoffery Hinton

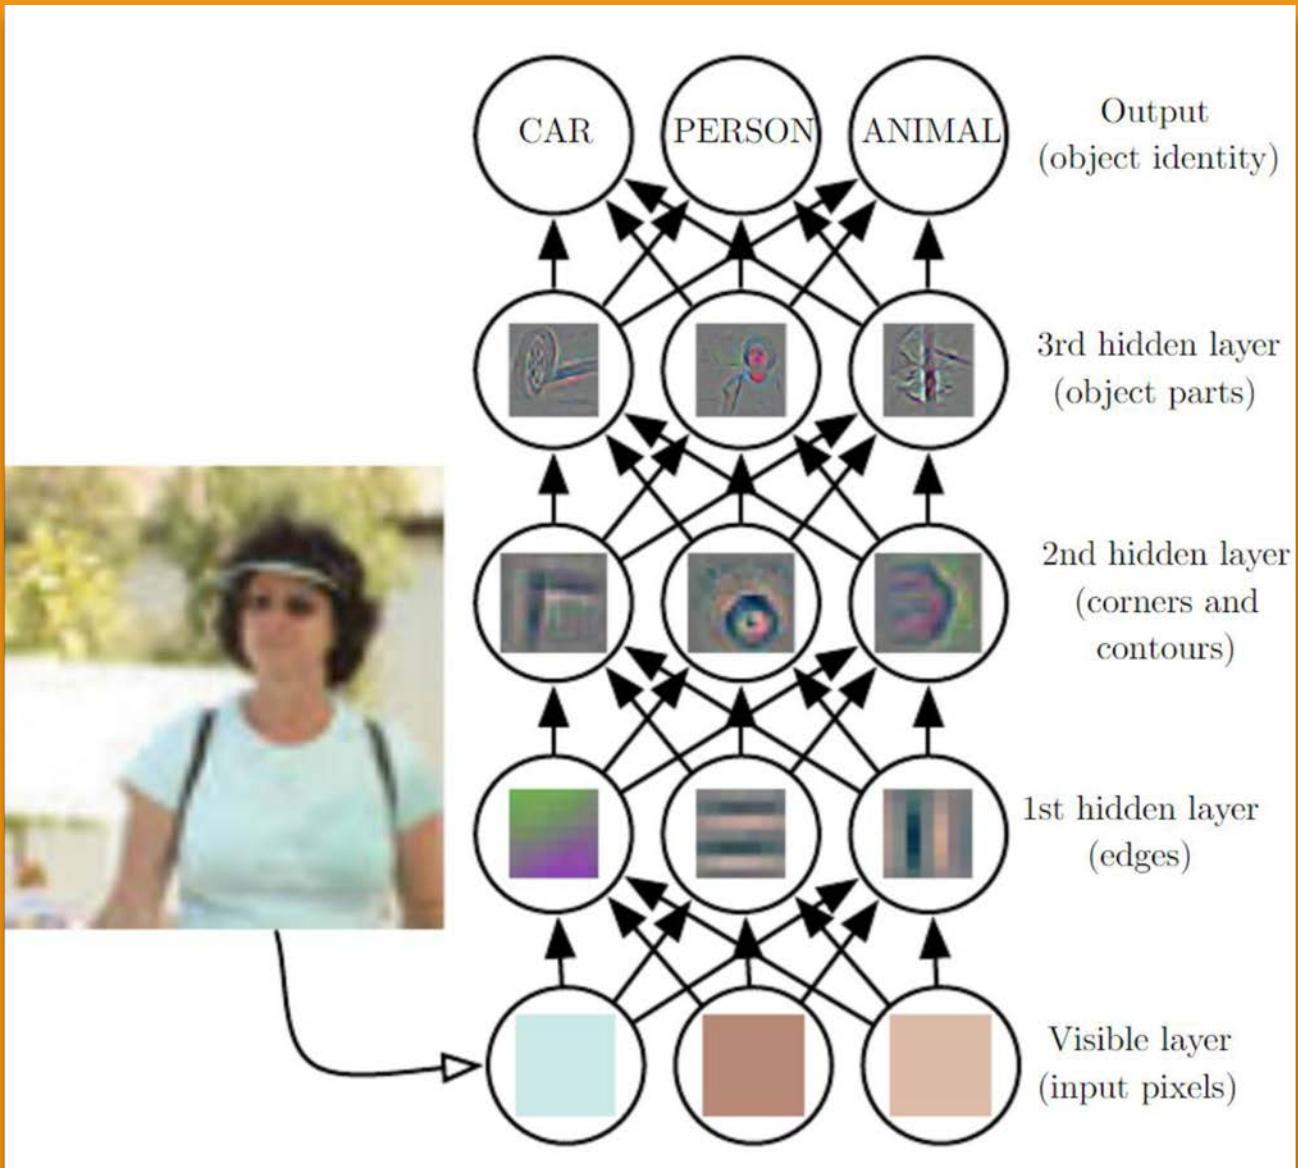


From: Gartner (August 2018)

# Deep Learning is Representation Learning

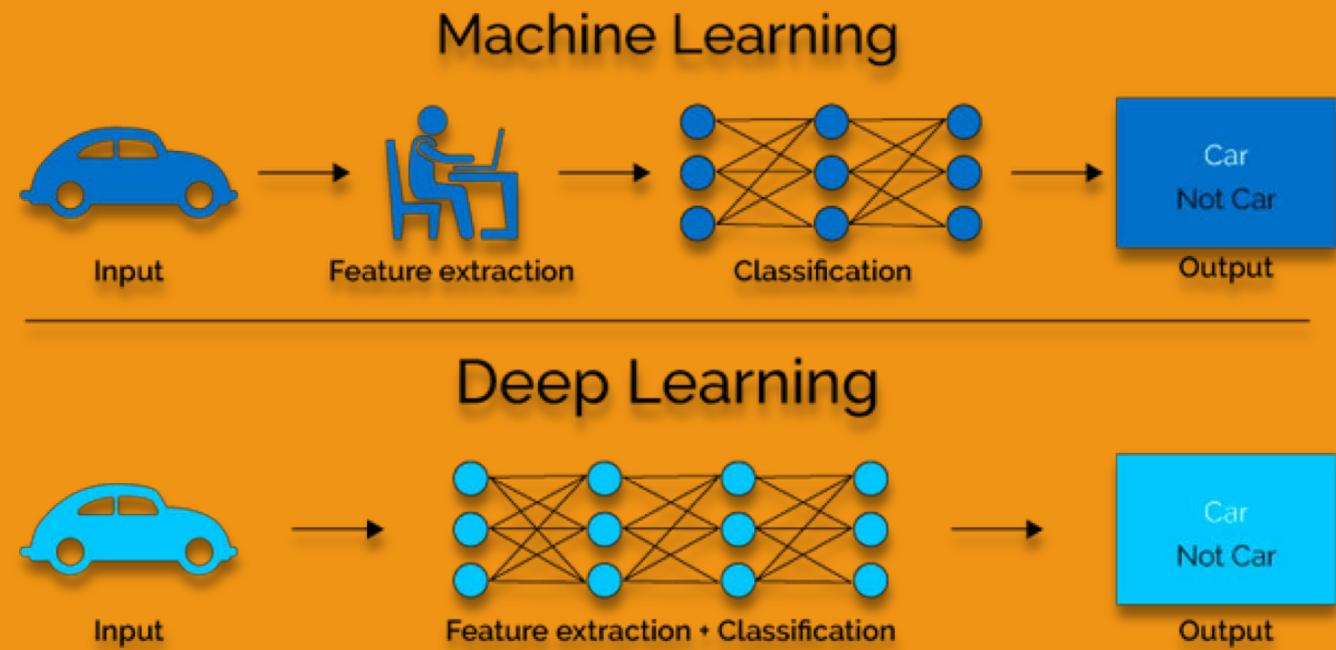


From: MIT Deep Learning (<https://deeplearning.mit.edu/>)



# Deep Learning: Automatic Feature Extraction

- More Data = Better Performance
- End-to-end Learning

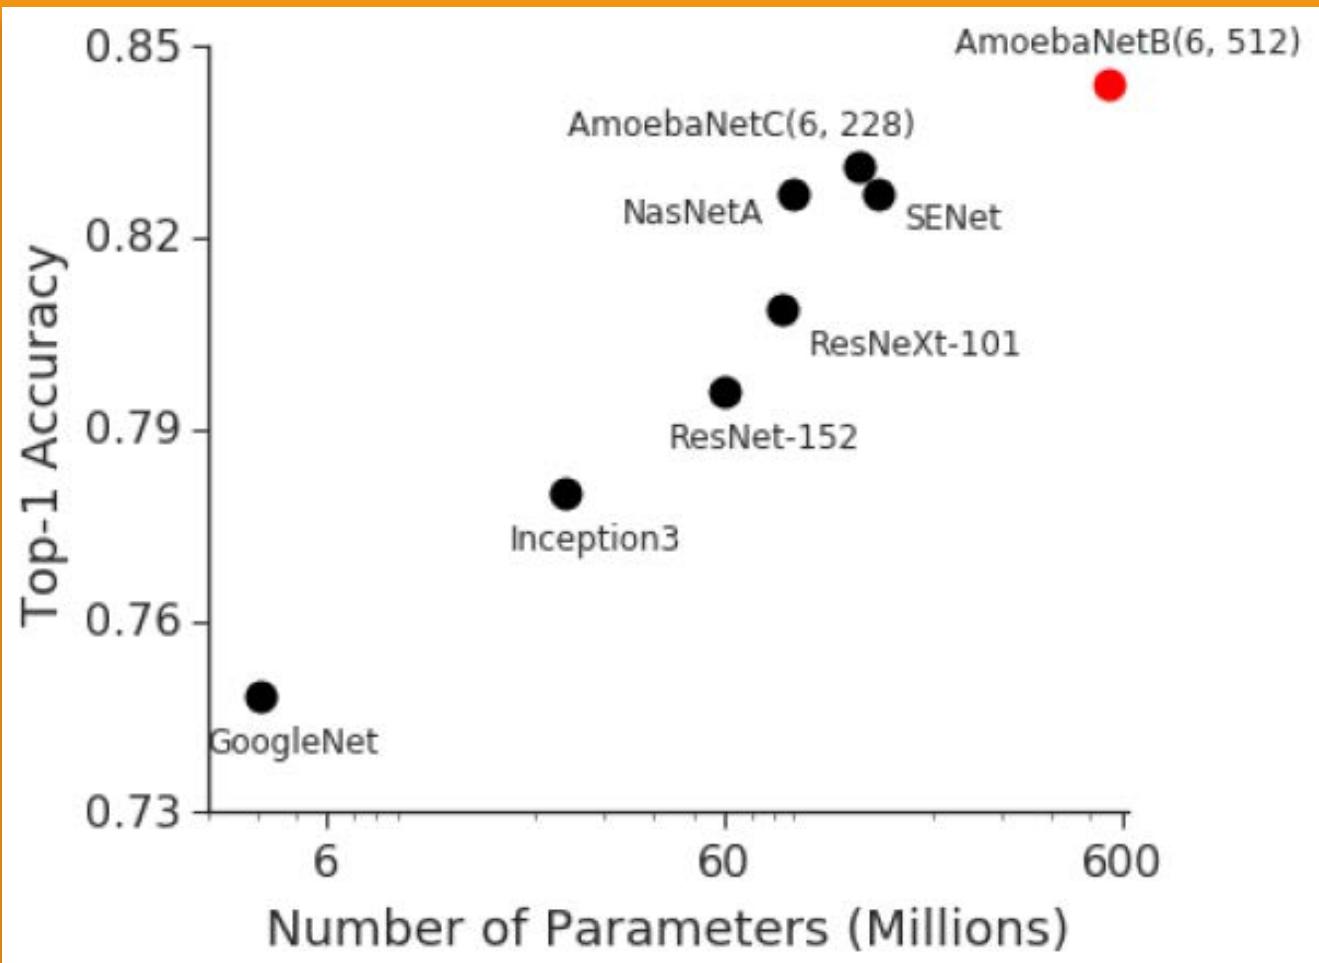


From: [codeutsava.in](http://codeutsava.in)

## Model Complexity

- More Data
- More Parameters
- Training Longer
- More Computation
- Better Results

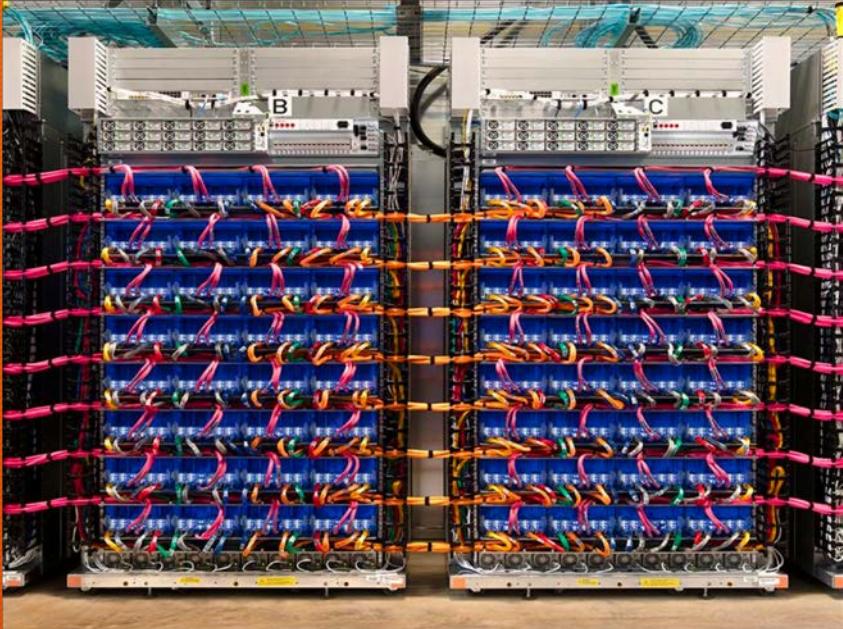
### ImageNet Object Detection Challenge



From: Introducing Gpipe, an Open Source Library for Efficiently  
Training Large-scale Neural Network Models (Google AI Blog)

# Hardware

Google Cloud TPU V2 Pod Alpha  
11.5 peta FLOPs



**FPGA**  
Programmable  
Hardware



**(GP)GPU**  
Data-Level Parallelism



**TPU**  
Custom ASIC specialized  
for machine learning



**CPU**  
Serial, General Purpose,  
Thread-level Parallelism



**Apple A12**

64-bit ARM 2.49GHz  
8-core Neural Engine



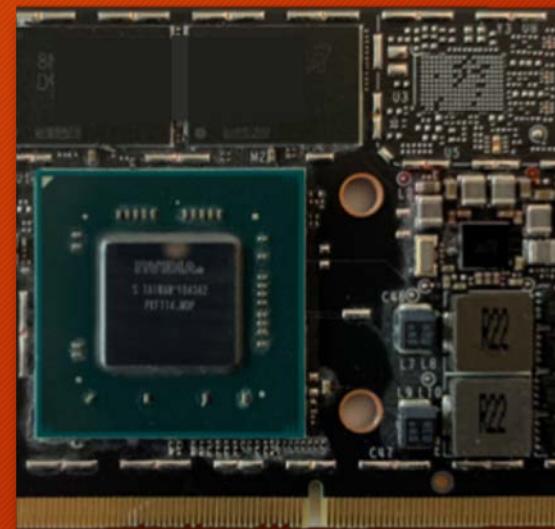
**Google Coral Dev**

NXP Quad Cortex-A53  
Google Edge TPU coprocessor



**Intel Neural Compute Stick**

USB 3.0 Stick  
Deep Learning Accelerator



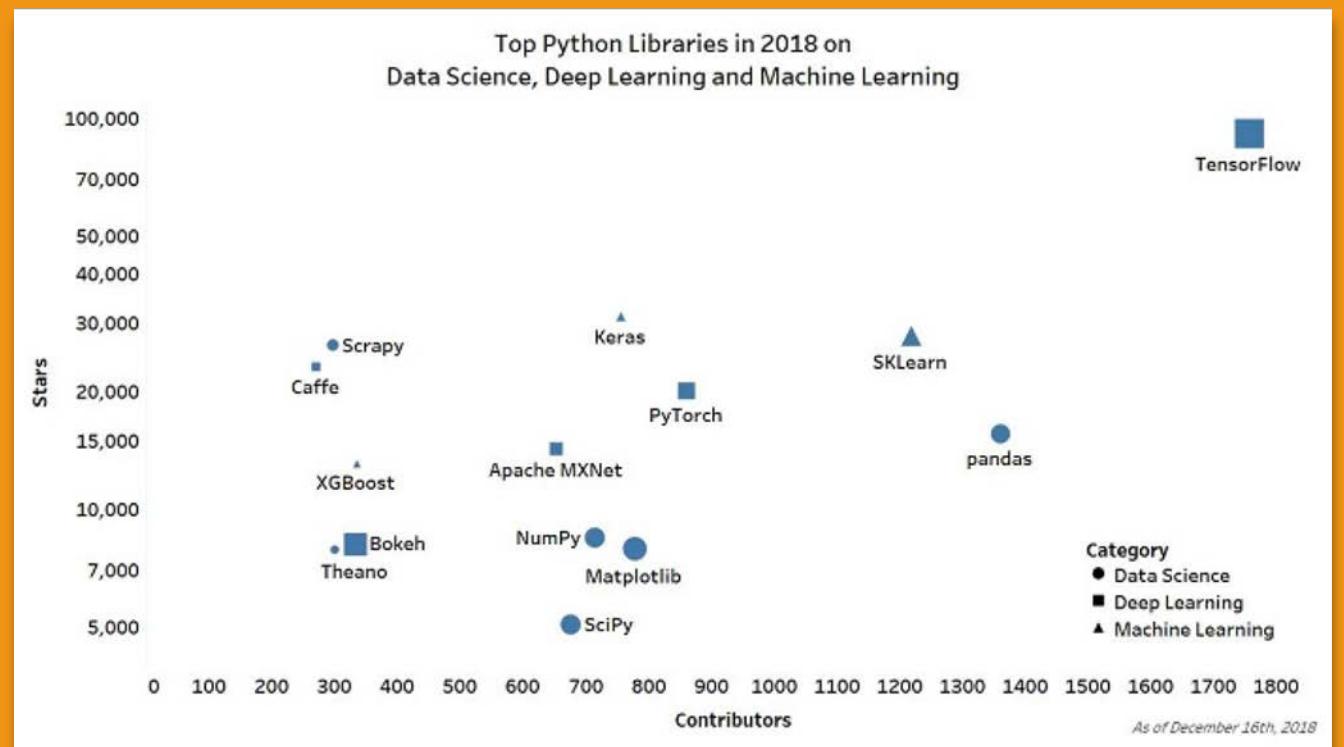
**Nvidia Jetson Nano**

Quad ARM A57  
Nvidia Maxwell, 128 CUDA cores

# Edge Devices

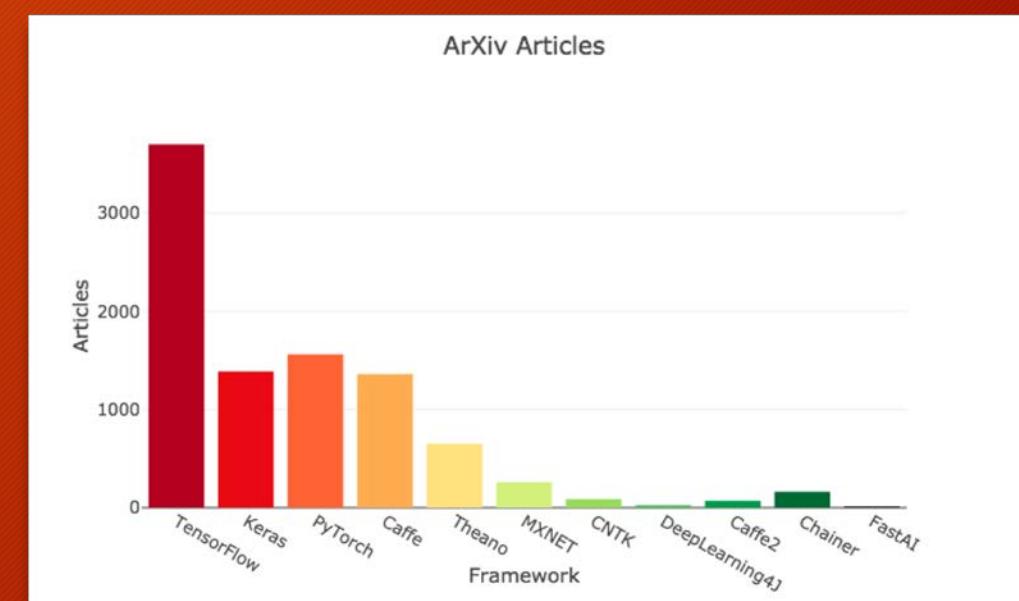
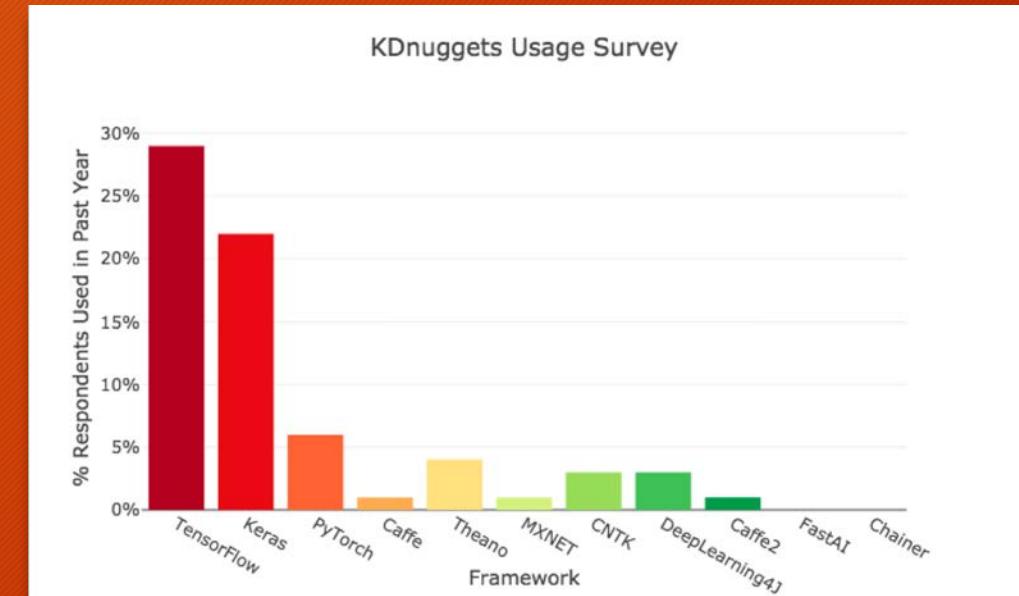
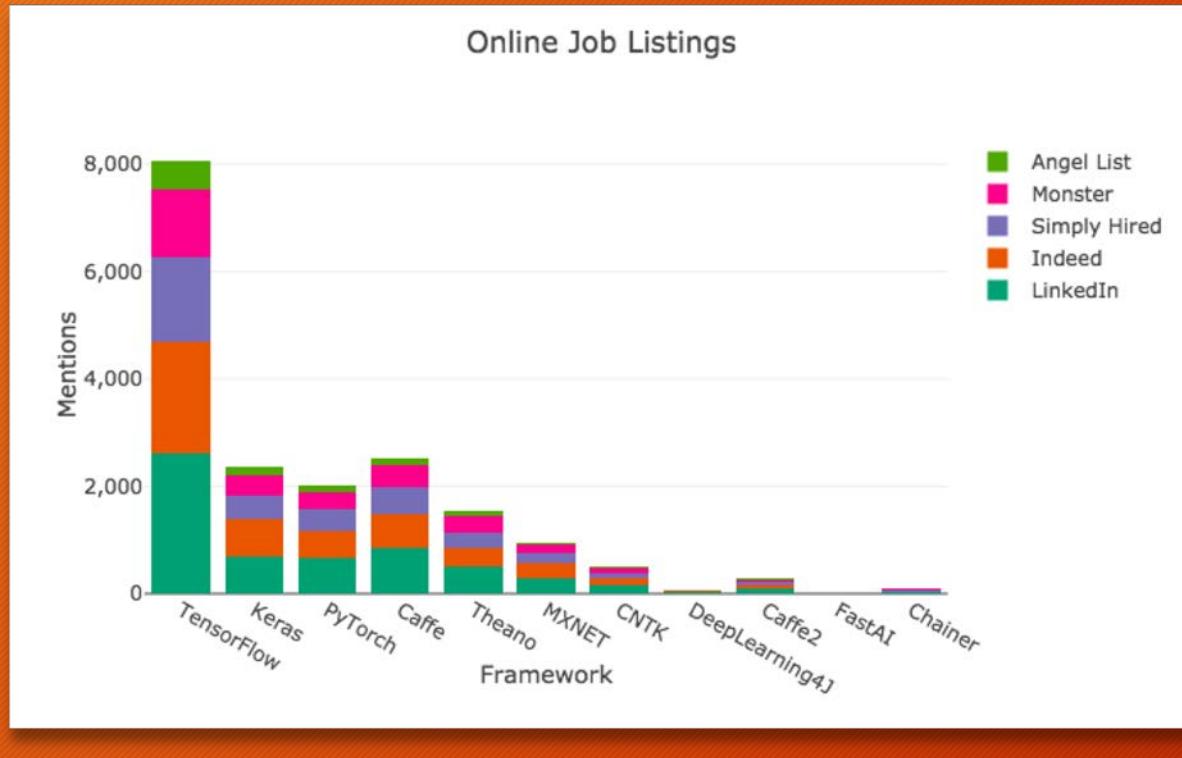
# Deep Learning Frameworks

- Tensorflow (Google)
- Keras (Google)
- pyTorch (Facebook)
- Caffe (UC, Berkeley)
- Theano (Universite de Montereal)
- MXNet (Apache)
- CNTK (Microsoft)
- DeepLearning4j (Eclipse)
- Chainer (Preferred Networks)
- Fast.ai



From: Top Python Libraries in 2018 in Data Science, Deep Learning, Machine Learning (KDnuggets)

# Framework Popularity



# TensorFlow in One Slide



## What is it:

Deep Learning Library  
Application Programming Interface  
(and more)



## Community:

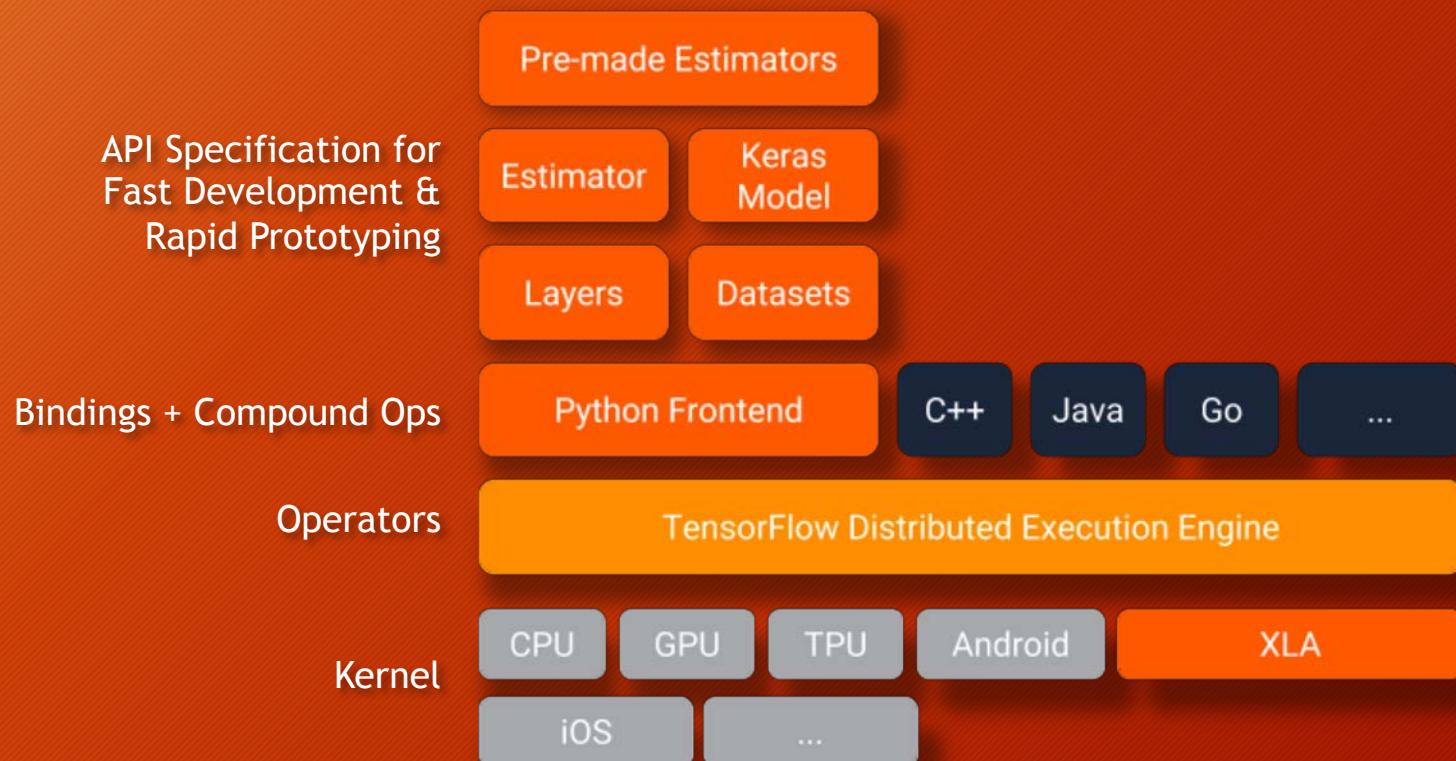
117,000+ Github Stars  
46,000 StackOverflow Questions  
[TensorFlow.org](http://TensorFlow.org)  
Blogs, Documentation, DevSummit,  
YouTube Talks



## Ecosystem:

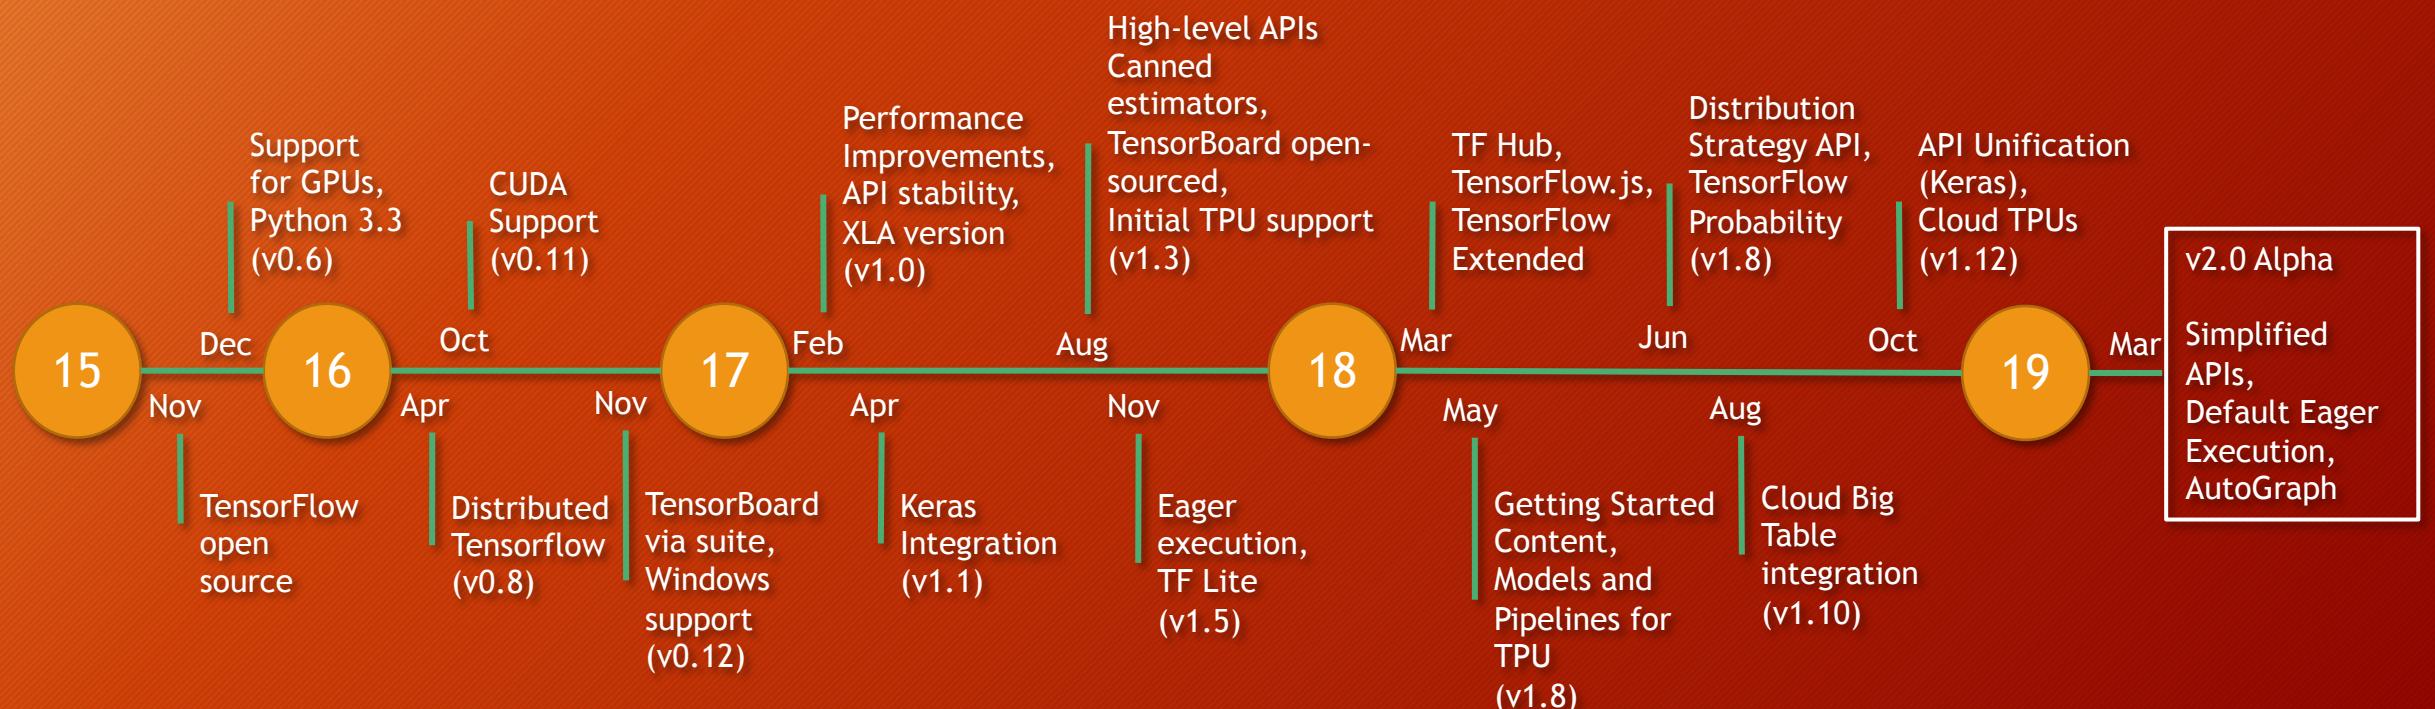
Keras: high-level API  
TensorFlow.js: in browser  
TensorFlow Lite: on the phone  
TPU: Optimized hardware  
TensorBoard: visualization  
TensorFlow Hub: Graph modules

# Tensorflow: Architecture



From: Introduction to TensorFlow Datasets and Estimators (Google Developers Blog)

# TensorFlow: Roadmap



# TF: Installation (Python)

- Installation with pip
  - With GPU support
  - Without GPU Support
- For best performance:
  - Compile from source (<https://github.com/tensorflow/tensorflow>)
- Nvidia CUDA Support
  - CUDA 10
  - CUDNN 7.4

# Python Environment Setup

- MiniConda 3 - Python 3.7.1
  - TensorFlow / TensorFlow-GPU
  - Scikit-Learn
  - Matplotlib
- Nvidia GPU with CUDA 10 (optional)
- Python IDE
  - Visual Studio Code
  - PyCharm (>=2019.1)

# TensorFlow: Getting Started

# What is Tensor?

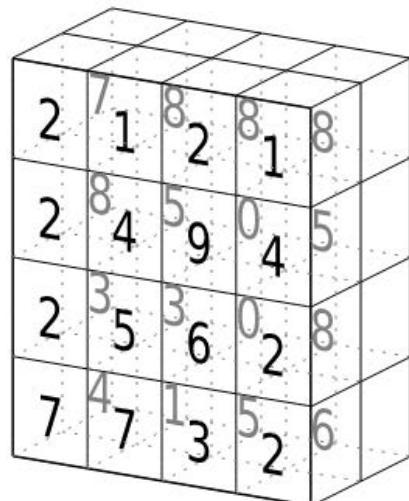
- Multi-dimensional Array of Numbers
- This means:
  - A scalar (number) is a tensor (0-d)
  - A vector is a tensor (1-d)
  - A matrix is a tensor ( $>=2$ -d)

't'
'e'
'n'
's'
'o'
'r'

tensor of dimensions [6]  
(vector of dimension 6)

3	1	4	1
5	9	2	6
5	3	5	8
9	7	9	3
2	3	8	4
6	2	6	4

tensor of dimensions [6,4]  
(matrix 6 by 4)



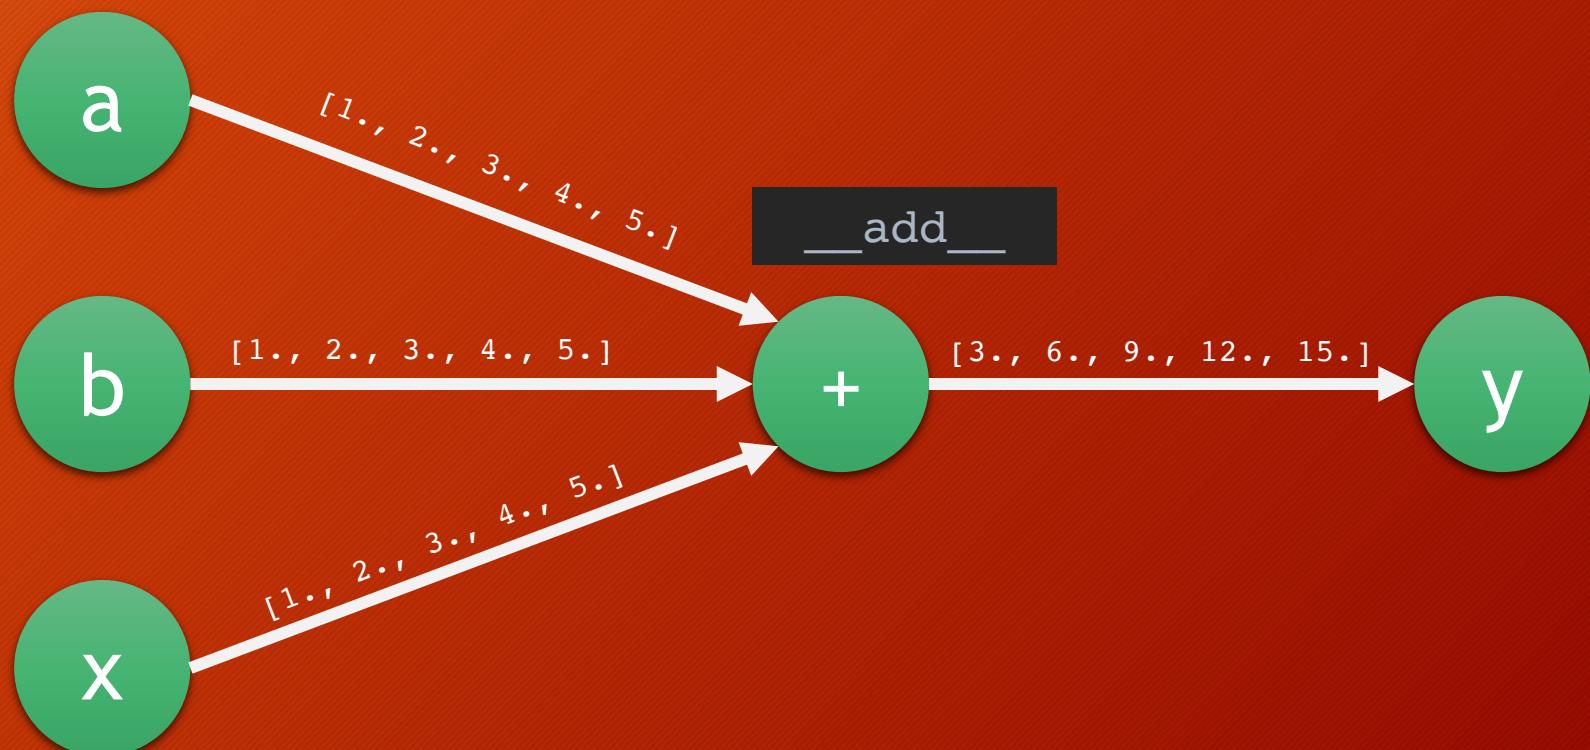
tensor of dimensions [4,4,2]

# A Simple Graph

```
tf.constant  
[1, 2, 3, 4, 5]
```

```
tf.Variable  
[1, 2, 3, 4, 5]
```

```
array  
[1, 2, 3, 4, 5]
```



# TF: Symbolic Computation

Build Graph

```
import tensorflow as tf

"""
Build Graph
"""

a = tf.constant(value=[1., 2., 3., 4., 5.], name="a")
b = tf.Variable(initial_value=[1., 2., 3., 4., 5.], name="b")

@tf.function
def SimpleAdder(x):
    return a + b + x

"""
Execute Graph
"""

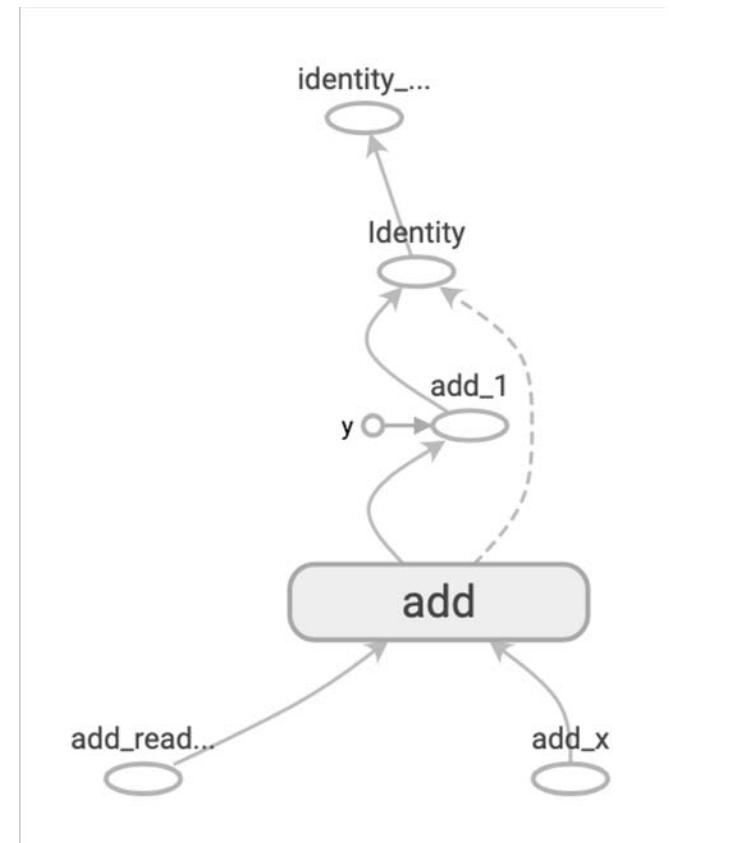
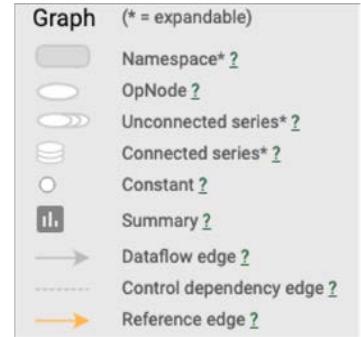
print(SimpleAdder([1., 2., 3., 4., 5.]))
```

Execute Graph

# TF: Data Flow Graph

- Nodes:
  - Operators, constants
- Edges:
  - Tensors

```
@tf.function
def SimpleAdder(x):
    return a + b + x
```



# TF 2.0: Functions, not Sessions

- Basic Idea: Python functions as Graphs

```
@tf.function  
def SimpleAdder(x):  
    return a + b + x
```

- The graph is executed when the function is invoked

```
SimpleAdder([1., 2., 3., 4., 5.])
```

# TensorBoard: Visualization

- Create file writer

```
logdir = 'logs/func/%s' % stamp  
writer = tf.summary.create_file_writer(os.path.abspath(logdir))
```

- Turn trace on before executing a graph

```
tf.summary.trace_on(graph=True)
```

- Export the trace to log directory

```
tf.summary.trace_export(name="my_func_trace", step=0, profiler_outdir=logdir)
```

- Open TensorBoard

```
(ml2019) $ tensorboard --logdir ./logs/func/
```

```
TensorBoard 1.14.0a20190301 at http://Tinghuis-iMac.local:6006 (Press CTRL+C to quit)
```

# Example 1: Simple Addition

# TF: Why Graph?

- Save computation
  - Only runs the subgraphs that lead to the values you want to fetch
- Auto Differentiation
- Facilitate distributed computation
  - The work can be spreaded across multiple CPUs, GPUs, TPUs or other devices.

# TF: Constant

```
import tensorflow as tf  
  
a = tf.constant([2, 2], name="a")  
b = tf.constant([1., 2.], name="b")  
c = tf.constant([2, 2], dtype=tf.float32, name="c")
```

```
d = tf.multiply(a, tf.transpose(b))  
  
>>> InvalidArgumentError: cannot  
compute Mul as input #1(zero-based) was  
expected to be a int32 tensor but is a  
float tensor [Op:Mul]
```

```
tf.constant(  
    value,  
    dtype=None,  
    shape=None,  
    name='Const'  
)
```

```
d = tf.multiply(c, tf.transpose(b))  
print(d)  
>>> tf.Tensor([2. 4.], shape=(2,), dtype=float32)  
  
tf.print(d)  
>>> [2 4]
```

# TF: Other Constant API

- Filled with a specific value

```
tf.zeros(shape, dtype=dtypes.float32, name=None)
tf.ones(shape, dtype=dtypes.float32, name=None)
tf.zeros_like(tensor, dtype=None, name=None, optimize=True)
tf.ones_like(tensor, dtype=None, name=None, optimize=True)
tf.fill(dims, value, name=None)
```

- Random Constants

```
tf.random.normal
tf.random.uniform
tf.random.categorical
tf.random.truncated_normal
tf.random.shuffle
tf.random.poisson
```

- Sequence

```
tf.linspace(start, stop, num, name=None)
tf.range(start, limit=None, delta=1, dtype=None, name="range")
```

# TF: Datatypes

- Similar to numpy
- Use “dtype” whenever possible

- `tf.float16` : 16-bit half-precision floating-point.
- `tf.float32` : 32-bit single-precision floating-point.
- `tf.float64` : 64-bit double-precision floating-point.
- `tf.bfloat16` : 16-bit truncated floating-point.
- `tf.complex64` : 64-bit single-precision complex.
- `tf.complex128` : 128-bit double-precision complex.
- `tf.int8` : 8-bit signed integer.
- `tf.uint8` : 8-bit unsigned integer.
- `tf.uint16` : 16-bit unsigned integer.
- `tf.int16` : 16-bit signed integer.
- `tf.int32` : 32-bit signed integer.
- `tf.int64` : 64-bit signed integer.
- `tf.bool` : Boolean.
- `tf.string` : String.
- `tf.qint8` : Quantized 8-bit signed integer.
- `tf.quint8` : Quantized 8-bit unsigned integer.
- `tf.qint16` : Quantized 16-bit signed integer.
- `tf.quint16` : Quantized 16-bit unsigned integer.
- `tf.qint32` : Quantized 32-bit signed integer.
- `tf.resource` : Handle to a mutable resource.

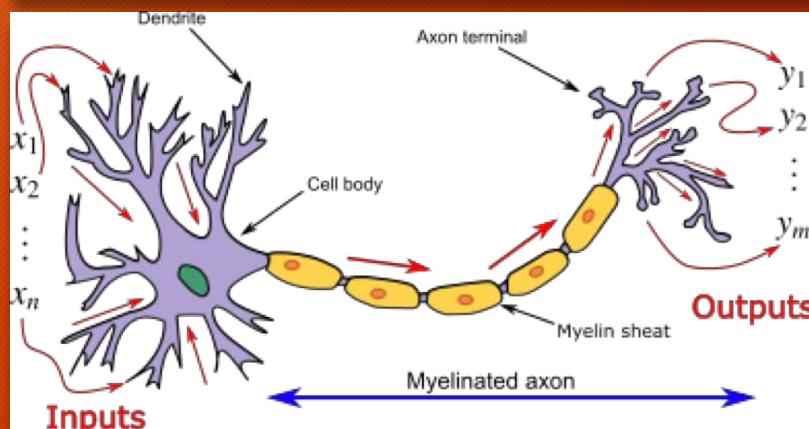
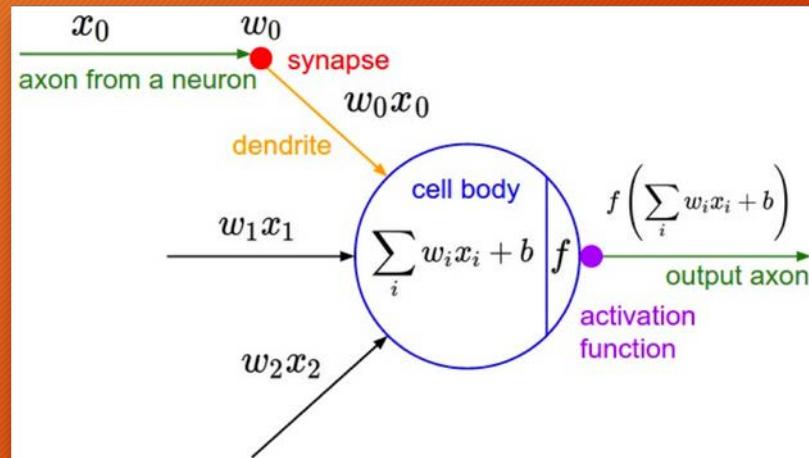
# What's wrong with Constants?

- Constants are stored in the graph definition
- This makes loading graphs expensive when constants are big
  
- Only use constants for primitive types
- Use variables or readers for more data that requires more memory

# TF: Operations

Category	Examples
Element-wise Math Ops	Add, Sub, Mul, Div, Exp, Log, Greater, Less, Equal...
Array Ops	Concat, Slice, Split, Constant, Rank, Shape, Shuffle...
Matrix Ops	MatMul, MatrixInverse, MatrixDeterminant, ...
Stateful Ops	Variable, Assign, ...
Neural Network Ops	SoftMax, Sigmoid, ReLU, Conv1D, Conv2D, MaxPool, ...
Common Loss Ops	CrossEntropy, L1/2 Regularization, RMS, ...
Optimization Ops	SGD, Adam, AdaDelta, Adagrad, RMSprop, ...
Checkpointing Operations	Save, Restore
Python Language	if, for, while, continue, break, ... (2.0)

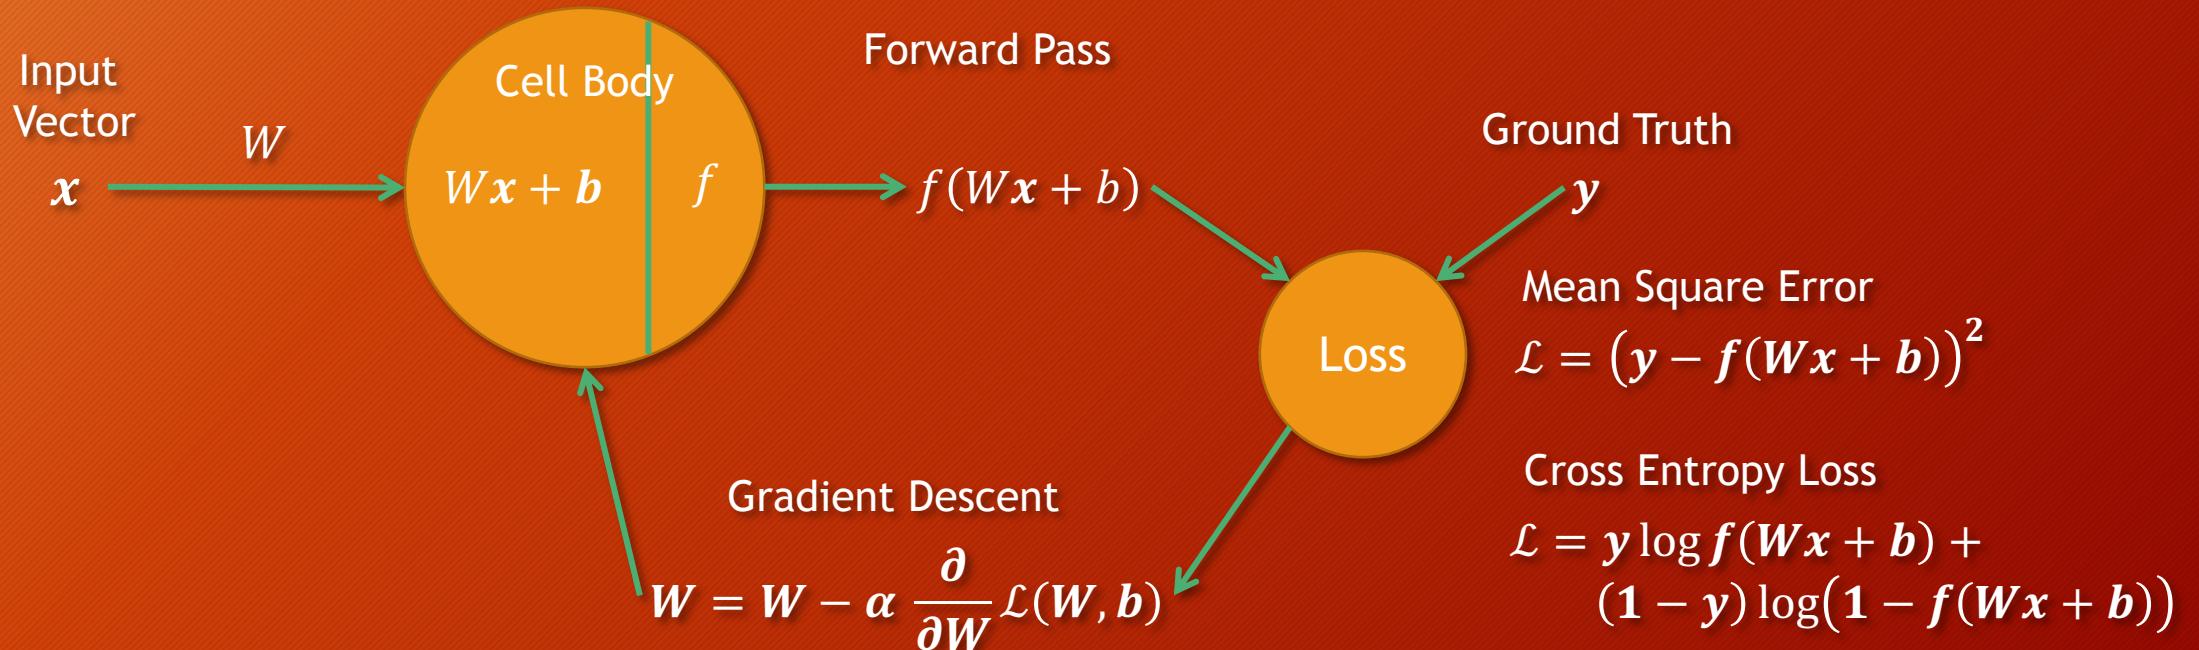
# Neural Network: Biological Inspiration for Computation



**Artificial Neuron:**  
computational building block for the “neural network”

**Neuron:**  
Computational building block for the brain

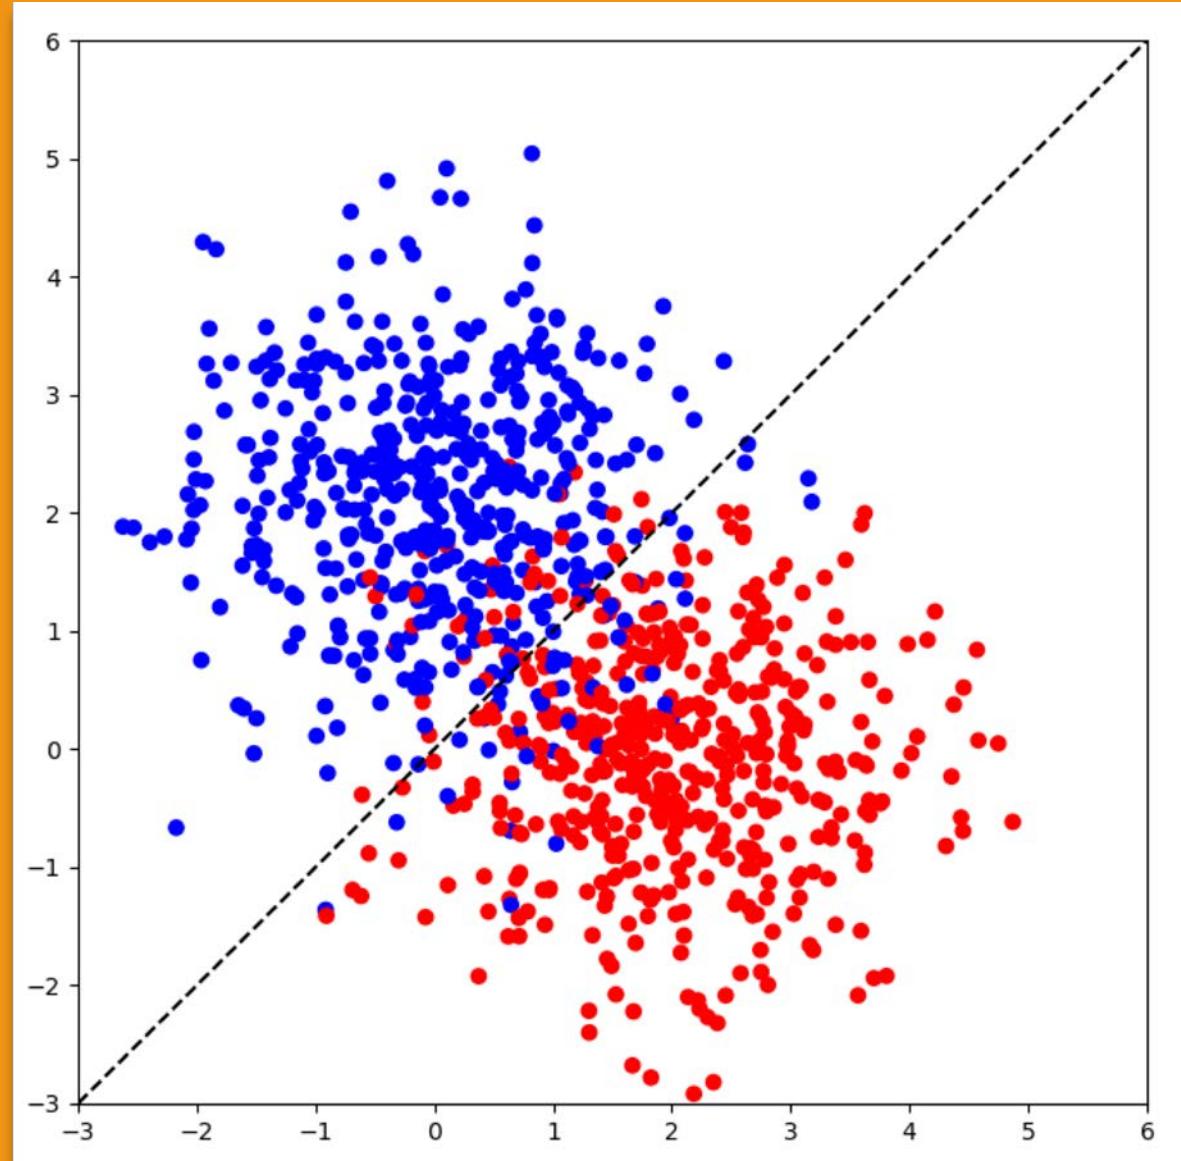
# Gradient Descent



## Example 2: Linear Regression (Custom Implementation)

## Dataset

- Synthetic Dataset
- Training set: 1,000
- Test set: 100
- Linear Classification



# Craft the model

```
w = tf.Variable(initial_value=np.random.normal(loc=0., scale=1., size=(2, 1)))  
b = tf.Variable(initial_value=np.random.normal((1,)))
```

```
@tf.function  
def linearModel(x):  
    return tf.matmul(x, w) + b
```

X: (batch\_size, 2)    W: (2, 1)    b: (1,)

# Training 1: Create Loss

```
def loss(y_true, logits):
    return tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(
        labels=y_true, logits=logits, name="cross_entropy_loss"
    ))
```

# Training #2: Calculate Gradient

```
@tf.function
def train(x, y, learning_rate):
    with tf.GradientTape() as t:
        logits = linearModel(x)
        current_loss = loss(y, logits)
    dw, db = t.gradient(current_loss, [w, b])
    w.assign_sub(learning_rate * dw)
    b.assign_sub(learning_rate * db)
    return current_loss
```

Compute gradient w.r.t.  
 $W$  and  $b$

# Training #3: Update Variable

```
@tf.function
def train(x, y, learning_rate):
    with tf.GradientTape() as t:
        logits = linearModel(x)
        current_loss = loss(y, logits)
    dw, db = t.gradient(current_loss, [w, b])
    w.assign_sub(learning_rate * dw)
    b.assign_sub(learning_rate * db)
    return current_loss
```

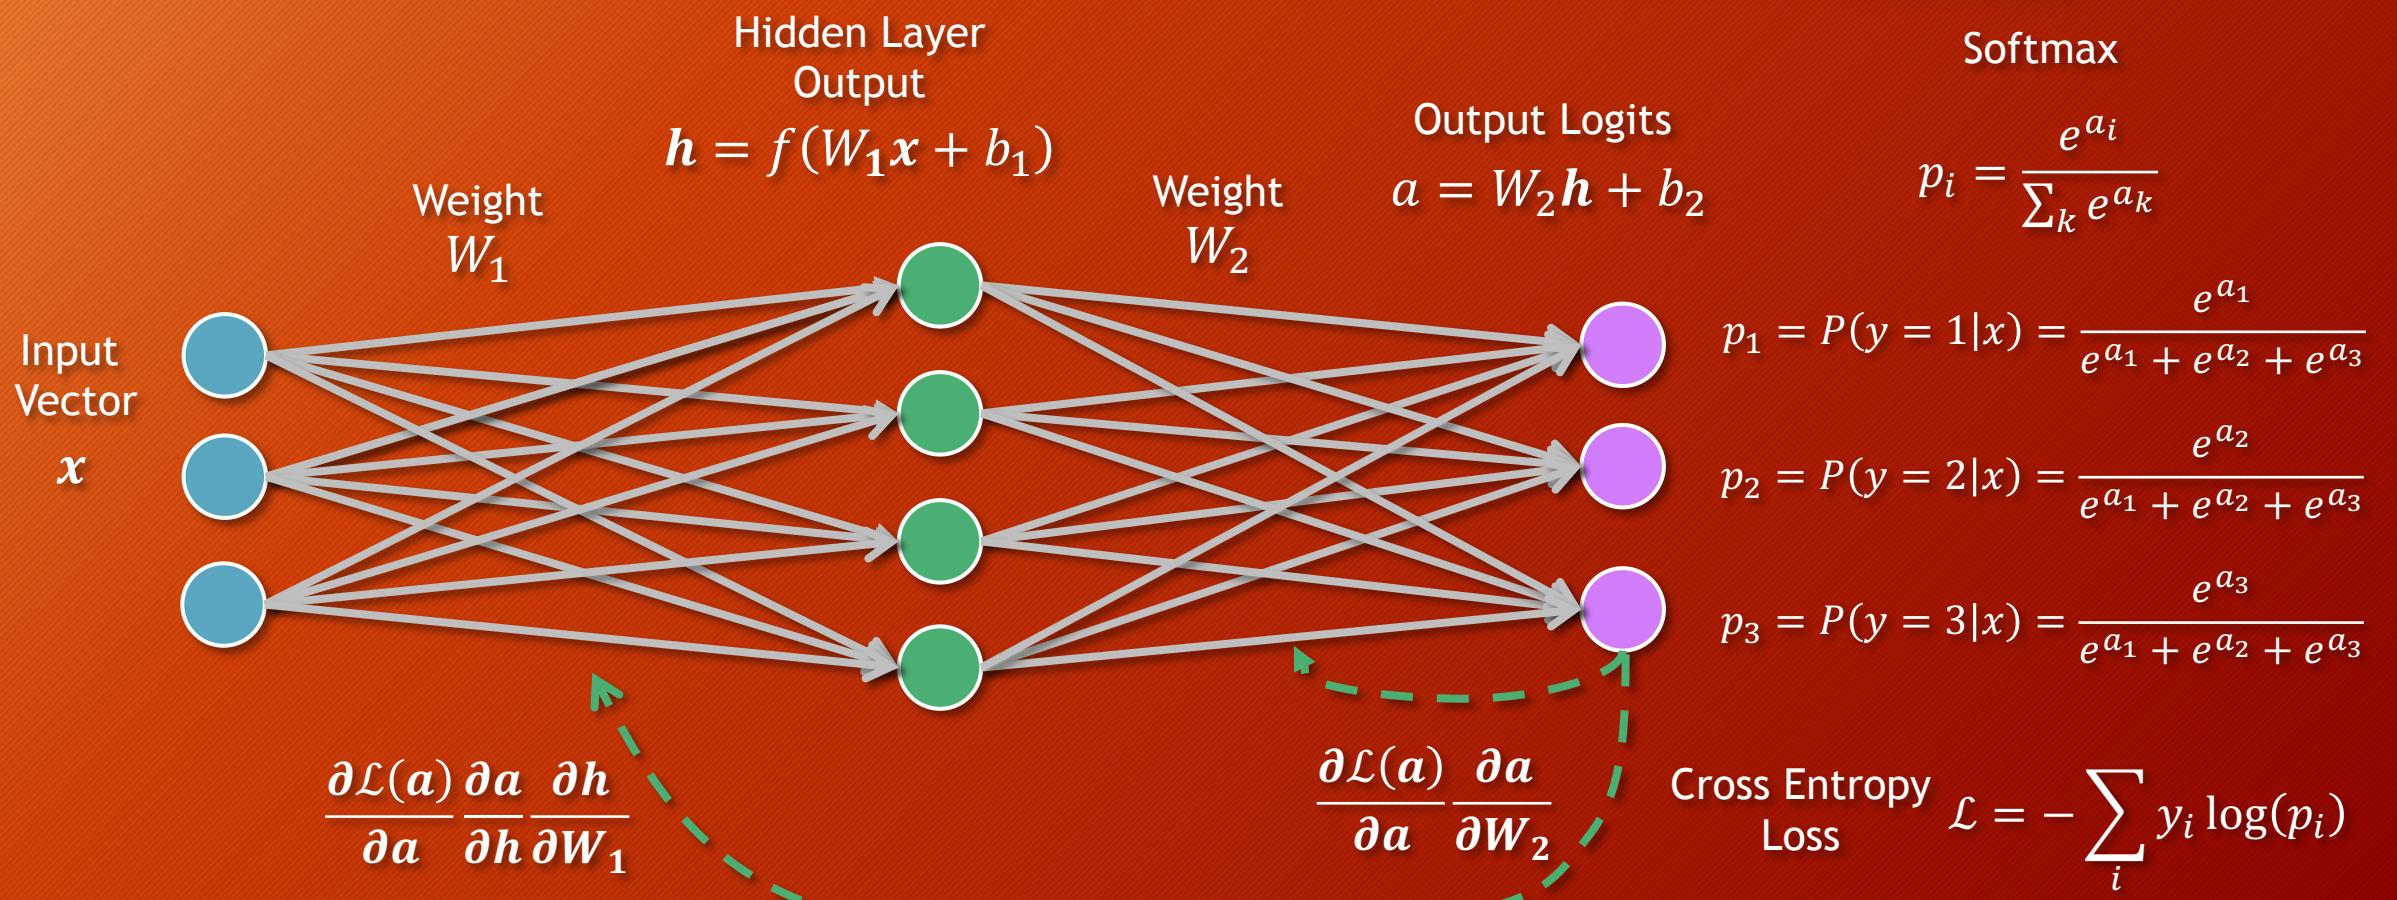
$$W := W - lr \frac{\partial \mathcal{L}}{\partial W}$$

# Training #4: Call training function

```
epochs = range(100)
for epoch in epochs:
    current_loss = train(
        x_train, y_train.reshape((1000, 1)).astype(dtype=np.float),
        learning_rate=1)
    print("Epoch %02d: loss %.8f" % (epoch, current_loss))
```

# Example 3: MNIST with MLP (Keras API)

# Multi-Layer Perceptron



# Multi-Layer Perceptron: Simple Implementation

1  
• Load Dataset

2  
• Build Forward Model

3  
• Loss Function and Optimizer

4  
• Training the Model

5  
• Evaluation

1

2

3

4

5

```
# For auto-completion to work
from sklearn.metrics import classification_report
import tensorflow.python.keras as keras

if __name__ == "__main__":
    (x_train, y_train), (x_test, y_test) =
        keras.datasets.mnist.load_data()
    # Normalize image pixel intensity
    x_train = x_train / 255
    x_test = x_test / 255

    # Build models with keras
    model = keras.Sequential([
        keras.layers.Flatten(input_shape=(28, 28)),
        keras.layers.Dense(128, activation='relu'),
        keras.layers.Dense(10, activation='softmax')
    ])

    model.compile(
        optimizer='adam',
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy']
    )

    model.fit(
        x=x_train, y=y_train, epochs=5
    )

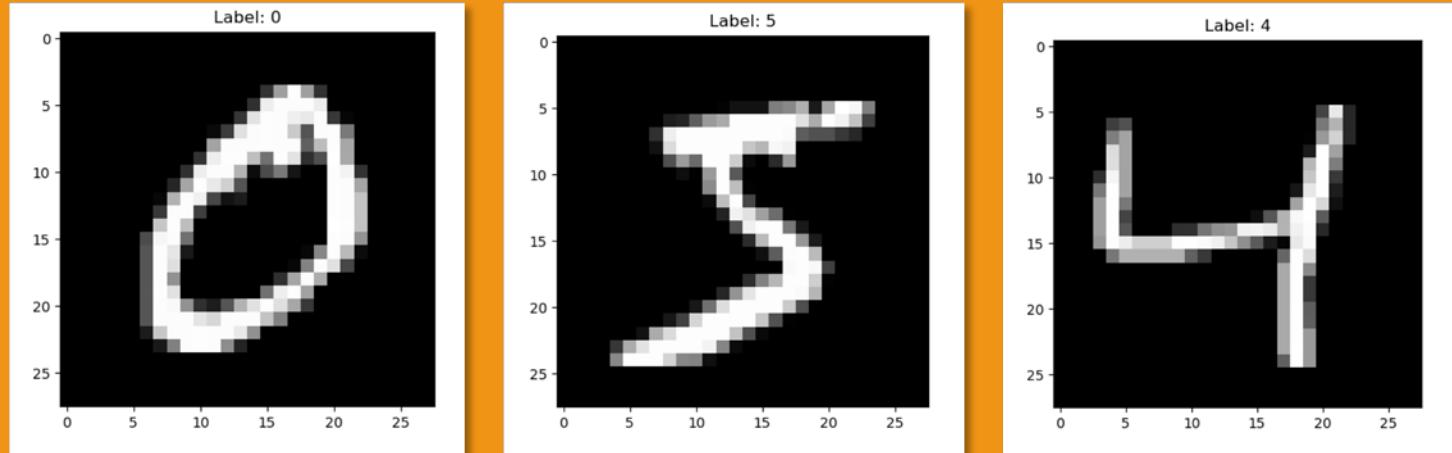
    y_pred = model.predict_classes(x=x_test)
    print(classification_report(y_test, y_pred, target_names=[
        '%d' % i for i in range(10)
    ]))
```

# MNIST Data

- Each image is 28x28 array
- Training Set: 60,000
- Test Set: 10,000

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()

for i in range(3):
    plt.figure()
    img = x_train[i, :, :]
    plt.title("Label: %d" % y_train[i])
    plt.imshow(img, cmap="gray")
    plt.savefig("digit_%d.png" % i)
    plt.show()
```



# How to learn TF?

- Take a Course
  - MIT: Deep Learning
  - Stanford: Introduction to TensorFlow
- Read Books
- Other's Codes
- Code to learn
  - Google Colaboratory

# TF: References

- [TensorFlow 2.0 Alpha Tutorial/Guide](#)
- [TensorFlow API 2.0](#)
- [TensorBoard For TF 2.0](#)
- [TensorFlow Blog](#)
- [TensorFlow Dev Summit](#)
- [TensorFlow RFCs](#)